

Lecture 4

Relational Model

Data Models and Database Design

- Think “logically” when design a database, need some kind of framework to design the database
- Like designing a data structure in some programming language, a data model is like a type system, but is abstract
- Organize the data into tables in the relational data model
- Initially no need to worry about how these tables are implemented

Relational Model Concepts

- The relational model of data is based on the concept of a **mathematical relation**
- A **relation** is a mathematical concept based on set theory
- The strength of the relational approach to data management comes from the formal foundation provided by the theory of relations

Relational Model Concepts

- The model was first proposed by Dr. E.F. Codd of IBM in 1970 in the following paper: "A Relational Model for Large Shared Data Banks," Communications of the ACM, June 1970

The above paper caused a major revolution in the field of Database management and earned Ted Codd the coveted Turing Award

What is a Relational Database?

Routes

| <u>RIId</u> | <u>RName</u> | <u>Grade</u> | <u>Rating</u> | <u>Height</u> |
|-------------|--------------|--------------|---------------|---------------|
| 1 | Last Tango | II | 12 | 100 |
| 2 | Garden Path | I | 2 | 60 |
| 3 | The Sluice | I | 8 | 60 |
| 4 | Picnic | III | 3 | 400 |

Climbers

| <u>CIId</u> | <u>Cname</u> | <u>Skill</u> | <u>Age</u> |
|-------------|--------------|--------------|------------|
| 123 | Edmund | EXP | 80 |
| 214 | Arnold | BEG | 25 |
| 313 | Bridget | EXP | 33 |
| 212 | James | MED | 27 |

Climbs

| <u>CIId</u> | <u>RIId</u> | <u>Date</u> | <u>Duration</u> |
|-------------|-------------|-------------|-----------------|
| 123 | 1 | 10/10/88 | 5 |
| 123 | 3 | 11/08/87 | 1 |
| 313 | 1 | 12/08/89 | 5 |
| 214 | 2 | 08/07/92 | 2 |
| 313 | 3 | 06/07/94 | 3 |

Can two columns in the same relation have the same name?

Yes

No

Total Results: 0

Can two columns in the same relation have the same name?

Yes

No

Can two columns in the same relation have the same name?

Yes

No

Why is the Database Like This?

- Each **route** has an *id*, a *name*, a *grade* (an estimate of the time needed), a *rating* (how difficult it is), and a *height*
- Each **climber** has an *id*, a *name*, a *skill* level and an *age*
- A **climb** records who climbed what route on what *date* and how long it took (*duration*)
- The data values in these tables are all “simple”

Describing Relations

- Relations are described by a schema which can be expressed in various ways
- A database schema is usually expressed in a data definition language (DDL)

```
Routes (RId:int, RName:string, Grade:string,  
        Rating:int, Height:int)
```

```
Climbers (CId:int, CNname:string,  
          Skill:string, Age:int)
```

```
Climbs (CId:int, RId:int, Date:date,  
        Duration:int)
```

Expressing Constraints

- In SQL, constraints are defined as follows:

```
CREATE TABLE Climbers
(CId INTEGER,
 CName CHAR(20),
 Skill CHAR(4),
 Age INTEGER,
PRIMARY KEY (CId),
UNIQUE (CName, Age));

CREATE TABLE Climbs
(CId INTEGER,
 RId INTEGER,
 Date DATE,
 Duration INTEGER,
PRIMARY KEY (CId, RId, Date),
FOREIGN KEY (CId) REFERENCES
    Climbers,
FOREIGN KEY (RId) REFERENCES
    Routes);
```

Informal Definition

- RELATION: A table of values
 - A relation is a set of rows
 - A relation is alternatively a set of columns
 - Each row represents a fact that corresponds to a real-world **entity** or **relationship**
 - Each row has an item or set of items that uniquely identifies that row in the table
 - Sometimes row-ids or sequential numbers are assigned to identify the rows in the table
 - Each column is called by its column name or column header or attribute name

Relational Model Terminology

- Table = relation
- Column headers = attributes
- Row = tuple
- Possible values of each attribute = domain
 - E.g., the domain of **CName** is string, the domain of **Rating** is real
- Relation schema = relation name + attributes + other structure information
 - E.g., keys, other constraints
- Relation instance is the current set of rows for a relation schema
- Database schema = collection of relation schemas

Schema Definition

- The **schema** of a relation: $R(A_1, A_2, \dots A_n)$
- Relation schema R is defined over attributes $A_1, A_2, \dots A_n$
 - R : Name
 - Degree: # of attributes n
 - Cardinality: # of tuples (rows)

Example

CUSTOMER (Cust-id, Cust-name, Address, Phone#)

- CUSTOMER is a relation defined over four **attributes** Cust-id, Cust-name, Address, Phone#
- Each attribute has a **domain** or a set of valid values
 - E.g., the domain of Cust-id could be 6 digit numbers
- A **tuple** is an ordered set of values
 - <632895, "John Smith", "101 Main St. Atlanta, GA 30332", "(404) 894-2000">
- A relation may be regarded as a set of tuples (rows)

Formal Definition

- The relation is a subset of the Cartesian product of the domains
 - Domain is used in a specific role conveyed by the attribute name
 - E.g., attribute Cust-name is defined over the domain of strings of 26 characters. The role these strings play in the CUSTOMER relation is that of the name of customers
- Formally,
Given $R(A_1, A_2, \dots, A_n)$
 $r(R) \subseteq \text{dom}(A_1) \times \text{dom}(A_2) \times \dots \times \text{dom}(A_n)$
R: relation schema (**intension**)
 $r(R)$: a specific "value" or population of R (**extension**)

Example

- $D_1 = \{0, 1\}$
- $D_2 = \{a, b, c\}$
- $r(R) \subseteq D_1 \times D_2$
- For example: $r(R) = \{ \langle 0, a \rangle, \langle 0, b \rangle, \langle 1, c \rangle \}$
is one possible “state” or “population” or
“extension” r of the relation R , defined over
domains D_1 and D_2

Definition Summary

| <u>Informal Terms</u> | <u>Formal Terms</u> |
|-----------------------|----------------------|
| Table | Relation |
| Column name | Attribute |
| Row | Tuple |
| Values in a column | Domain |
| Table Definition | Schema of a Relation |
| Populated Table | Extension |

Example

The diagram illustrates a database relation table. The table is labeled 'STUDENT' in the first column. The columns are labeled 'Name', 'SSN', 'HomePhone', 'Address', 'OfficePhone', 'Age', and 'GPA'. The rows represent individual students. Annotations include: 'Relation name' pointing to the 'STUDENT' column header; 'Attributes' pointing to the column headers; and 'Tuples' pointing to the rows of data.

| STUDENT | Name | SSN | HomePhone | Address | OfficePhone | Age | GPA |
|---------|-----------------|-------------|-----------|----------------------|-------------|-----|------|
| | Benjamin Bayer | 305-61-2435 | 373-1616 | 2918 Bluebonnet Lane | null | 19 | 3.21 |
| | Katherine Ashly | 381-62-1245 | 375-4409 | 125 Kirby Road | null | 18 | 2.89 |
| | Dick Davidson | 422-11-2320 | null | 3452 Elgin Road | 749-1253 | 25 | 3.53 |
| | Charles Cooper | 489-22-1100 | 376-9821 | 265 Lark Lane | 749-6492 | 28 | 3.93 |
| | Barbara Benson | 533-69-1238 | 839-8461 | 7384 Fontana Lane | null | 19 | 3.25 |

More Relational Model Notations

- $t[A_i] = v_i$ (the value of attribute A_i for tuple t)
- $t[A_u, A_v, \dots, A_w]$: the tuple of t containing the values of attributes A_u, A_v, \dots, A_w

Characteristics of Relations

- **Ordering of tuples in a relation $r(R)$:**
 - Not ordered (like a set)
- **Ordering of attributes in a relation schema R**
 - Not important
 - However, the correspondence between the attributes in $R(A_1, A_2, \dots, A_n)$ and the values in $t = \langle v_1, v_2, \dots, v_n \rangle$ is maintained


Characteristics of Relations (Cont.)

- **Values in a tuple:**
 - A set of *atomic* values
 - Each value in the domain is indivisible
 - Composite and multi-valued attributes are not allowed
 - **null** value: unknown or inapplicable to certain tuples

Example

| STUDENT | Name | SSN | HomePhone | Address | OfficePhone | Age | GPA |
|---------|-----------------|-------------|-----------|----------------------|-------------|-----|------|
| | Dick Davidson | 422-11-2320 | null | 3452 Elgin Road | 749-1253 | 25 | 3.53 |
| | Barbara Benson | 533-69-1238 | 839-8461 | 7384 Fontana Lane | null | 19 | 3.25 |
| | Charles Cooper | 489-22-1100 | 376-9821 | 265 Lark Lane | 749-6492 | 28 | 3.93 |
| | Katherine Ashly | 381-62-1245 | 375-4409 | 125 Kirby Road | null | 18 | 2.89 |
| | Benjamin Bayer | 305-61-2435 | 373-1616 | 2918 Bluebonnet Lane | null | 19 | 3.21 |

Relational Data Model



Relational Model Constraints

- Constraints are *conditions* that must hold on **all** valid relation instances
- Three main constraints:
 1. **Key** constraints (single relation)
 2. **Entity integrity** constraints (single relation)
 3. **Referential integrity** constraints (two relations)

Key Constraints

- **Key constraints:** No two tuples can have the same value for the key
- **Superkey:** A set of attributes SK of R such that no two tuples *in any valid relation instance* $r(R)$ will have the same value for SK
 - For any distinct tuples t_1 and t_2 in $r(R)$, $t_1[SK] \neq t_2[SK]$
- **Candidate Key:** A "minimal" superkey
 - A superkey K such that removal of any attribute from K results in a set of attributes that is not a superkey

Example 1

CAR(State, Reg#, SerialNo, Make, Model, Year)

- Two candidate keys:
 - Key1 = {State, Reg#}
 - Key2 = {SerialNo}
 - Are Key1 and Key2 superkeys?
- {SerialNo, Make}: candidate key or superkey?
- **Primary key**: If a relation has several candidate keys, one is chosen arbitrarily to be the primary key
 - The primary key is underlined
 - Implies “not null”

Example 2

- The CAR relation, with two candidate keys: License-number and Engine_serial_number

CAR

| <u>License_number</u> | Engine_serial_number | Make | Model | Year |
|-----------------------|----------------------|------------|---------|------|
| Texas ABC-739 | A69352 | Ford | Mustang | 02 |
| Florida TVP-347 | B43696 | Oldsmobile | Cutlass | 05 |
| New York MPO-22 | X83554 | Oldsmobile | Delta | 01 |
| California 432-TFY | C43742 | Mercedes | 190-D | 99 |
| California RSK-629 | Y82935 | Toyota | Camry | 04 |
| Texas RSK-629 | U028365 | Jaguar | XJS | 04 |

Entity Integrity

- **Relational Database Schema:** A set of relation schemas that belong to the same database
 - $S = \{R_1, R_2, \dots, R_n\}$, S is the name of the **database**
- **Entity Integrity:** The PK (primary key) attributes of each relation schema R in S cannot have null values in any tuple of $r(R)$
 - $t[PK] \neq \text{null}$ for any tuple t in $r(R)$
 - Why?
- Other attributes of R may be similarly constrained to disallow null values, even though they are not members of the primary key

Referential Integrity

- If a tuple in R_1 refers to R_2 , it must refer to an **existing** tuple in R_2
- E.g., Dno in every Employee tuple must match Dnumber value of some tuple in the Department relation


Foreign Key Constraints

- A foreign key constraint mostly involves *two* relations
- Specify a *relationship* among tuples in two relations:
 - The **referencing relation** (R_1) and the **referenced relation** (R_2)
 - FK (**foreign key** attributes) of R_1 reference PK (primary key attributes) or candidate key of R_2
- Displayed in a relational database schema as a directed arc from $R_1.FK$ to $R_2.PK$ or $R_2.CK$

Foreign Key

- A set of attributes in R_1 is a foreign key of R_1 that references relation R_2 if it satisfies the following two rules:
 - The attributes in FK have the same domain(s) as the primary key (PK) attributes or candidate key of R_2
 - A value in the foreign key column (or columns) of the the **referencing relation** R_1 can be either:
 1. a value of an existing primary key or candidate key value in the **referenced relation** R_2
 2. is null
- In case (2), the FK in R_1 should not be a part of its own primary key

Example



- If no row exists in R2– violation of referential integrity
- Not all rows in R2 need to be referenced
- Value of a foreign key might not be specified
- Names of K1 and K2 need not to be the same

Does foreign key need to be unique?

Yes

No

Total Results: 0

Does foreign key need to be unique?

Yes

No

Does foreign key need to be unique?

Yes

No

Foreign Key (Cont.)

- Foreign key can refer to its own relation
- For example: SuperSSN is a FK that refers to SSN of Employee relation itself

Example 1

- Schema diagram for the COMPANY relational database schema (the primary keys are underlined)

| EMPLOYEE | | | | | | | | | |
|----------|-------|-------|------------|-------|---------|-----|--------|----------|-----|
| FNAME | MINIT | LNAME | <u>SSN</u> | BDATE | ADDRESS | SEX | SALARY | SUPERSSN | DNO |

| DEPARTMENT | | | |
|------------|----------------|--------|--------------|
| DNAME | <u>DNUMBER</u> | MGRSSN | MGRSTARTDATE |

| DEPT_LOCATIONS | |
|----------------|------------------|
| <u>DNUMBER</u> | <u>DLOCATION</u> |

| PROJECT | | | |
|---------|----------------|-----------|------|
| PNAME | <u>PNUMBER</u> | PLOCATION | DNUM |

| WORKS_ON | | |
|-------------|------------|-------|
| <u>ESSN</u> | <u>PNO</u> | HOURS |

| DEPENDENT | | | | |
|-------------|-----------------------|-----|-------|--------------|
| <u>ESSN</u> | <u>DEPENDENT_NAME</u> | SEX | BDATE | RELATIONSHIP |

Example 1 (Cont.)

- One possible relational database state corresponding to the COMPANY schema

EMPLOYEE

| Fname | Minit | Lname | Ssn | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|----------|-------|---------|-----------|------------|--------------------------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |
| Franklin | T | Wong | 333445555 | 1955-12-08 | 638 Voss, Houston, TX | M | 40000 | 888665555 | 5 |
| Alicia | J | Zelaya | 999887777 | 1968-01-19 | 3321 Castle, Spring, TX | F | 25000 | 987654321 | 4 |
| Jennifer | S | Wallace | 987654321 | 1941-06-20 | 291 Berry, Bellaire, TX | F | 43000 | 888665555 | 4 |
| Ramesh | K | Narayan | 666884444 | 1962-09-15 | 975 Fire Oak, Humble, TX | M | 38000 | 333445555 | 5 |
| Joyce | A | English | 453453453 | 1972-07-31 | 5631 Rice, Houston, TX | F | 25000 | 333445555 | 5 |
| Ahmad | V | Jabbar | 987987987 | 1969-03-29 | 980 Dallas, Houston, TX | M | 25000 | 987654321 | 4 |
| James | E | Borg | 888665555 | 1937-11-10 | 450 Stone, Houston, TX | M | 55000 | NULL | 1 |

DEPARTMENT

| Dname | Dnumber | Mgr_ssn | Mgr_start_date |
|----------------|---------|-----------|----------------|
| Research | 5 | 333445555 | 1988-05-22 |
| Administration | 4 | 987654321 | 1995-01-01 |
| Headquarters | 1 | 888665555 | 1981-06-19 |

DEPT_LOCATIONS

| Dnumber | Dlocation |
|---------|-----------|
| 1 | Houston |
| 4 | Stafford |
| 5 | Bellaire |
| 5 | Sugarland |
| 5 | Houston |

WORKS_ON

| Essn | Pno | Hours |
|-----------|-----|-------|
| 123456789 | 1 | 32.5 |
| 123456789 | 2 | 7.5 |
| 666884444 | 3 | 40.0 |
| 453453453 | 1 | 20.0 |
| 453453453 | 2 | 20.0 |
| 333445555 | 2 | 10.0 |
| 333445555 | 3 | 10.0 |
| 333445555 | 10 | 10.0 |
| 333445555 | 20 | 10.0 |
| 999887777 | 30 | 30.0 |
| 999887777 | 10 | 10.0 |
| 987987987 | 10 | 35.0 |
| 987987987 | 30 | 5.0 |
| 987654321 | 30 | 20.0 |
| 987654321 | 20 | 15.0 |
| 888665555 | 20 | NULL |

PROJECT

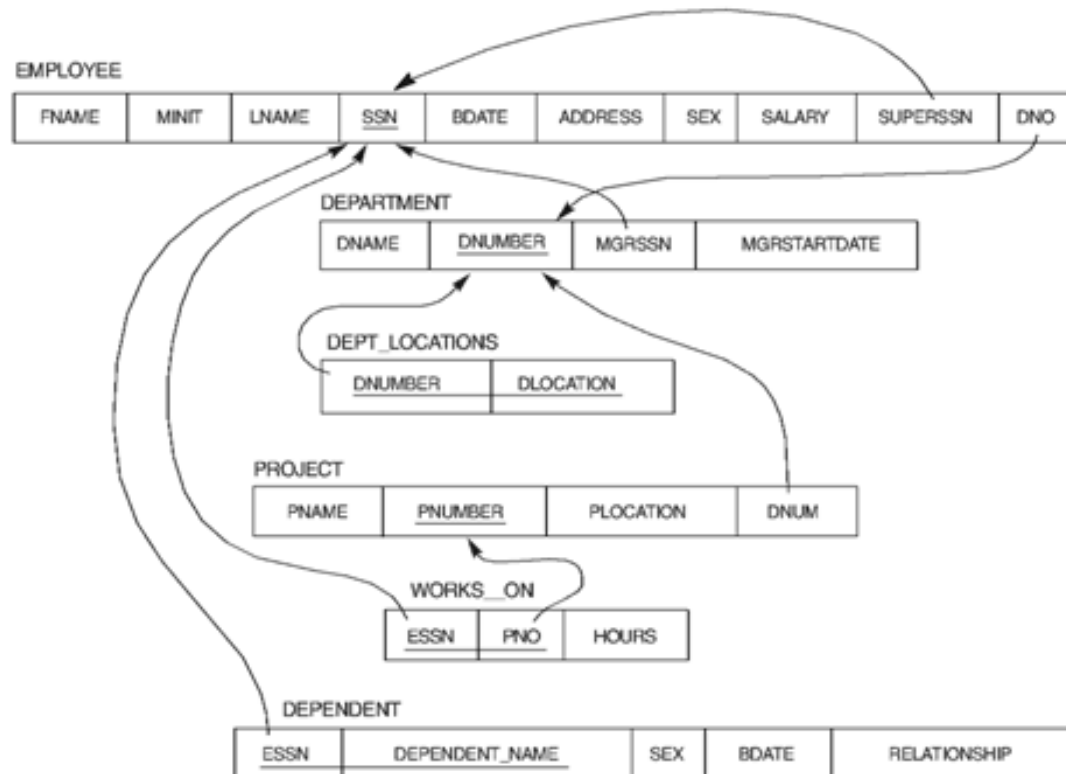
| Pname | Pnumber | Plocation | Dnum |
|-----------------|---------|-----------|------|
| ProductX | 1 | Bellaire | 5 |
| ProductY | 2 | Sugarland | 5 |
| ProductZ | 3 | Houston | 5 |
| Computerization | 10 | Stafford | 4 |
| Reorganization | 20 | Houston | 1 |
| Newbenefits | 30 | Stafford | 4 |

DEPENDENT

| Essn | Dependent_name | Sex | Bdate | Relationship |
|-----------|----------------|-----|------------|--------------|
| 333445555 | Alice | F | 1986-04-05 | Daughter |
| 333445555 | Theodore | M | 1983-10-25 | Son |
| 333445555 | Joy | F | 1958-05-03 | Spouse |
| 987654321 | Abner | M | 1942-02-28 | Spouse |
| 123456789 | Michael | M | 1988-01-04 | Son |
| 123456789 | Alice | F | 1988-12-30 | Daughter |
| 123456789 | Elizabeth | F | 1967-05-05 | Spouse |

Example 1 (Cont.)

- Foreign key constraints displayed on the COMPANY relational database schema diagram



Example 2

AIRPORT

| | | | |
|--------------------|------|------|-------|
| <u>airportcode</u> | name | city | state |
|--------------------|------|------|-------|

FLT-SCHEDULE

| | | | | | | | |
|-------------|---------|-------|------------------|-------|----------------|-------|-------|
| <u>flt#</u> | airline | dtime | from-airportcode | atime | to-airportcode | miles | price |
|-------------|---------|-------|------------------|-------|----------------|-------|-------|

FLT-WEEKDAY

| | |
|-------------|----------------|
| <u>flt#</u> | <u>weekday</u> |
|-------------|----------------|

FLT-INSTANCE

| | | | |
|-------------|-------------|--------|--------------|
| <u>flt#</u> | <u>date</u> | plane# | #avail-seats |
|-------------|-------------|--------|--------------|

AIRPLANE

| | | |
|---------------|------------|--------------|
| <u>plane#</u> | plane-type | total-#seats |
|---------------|------------|--------------|

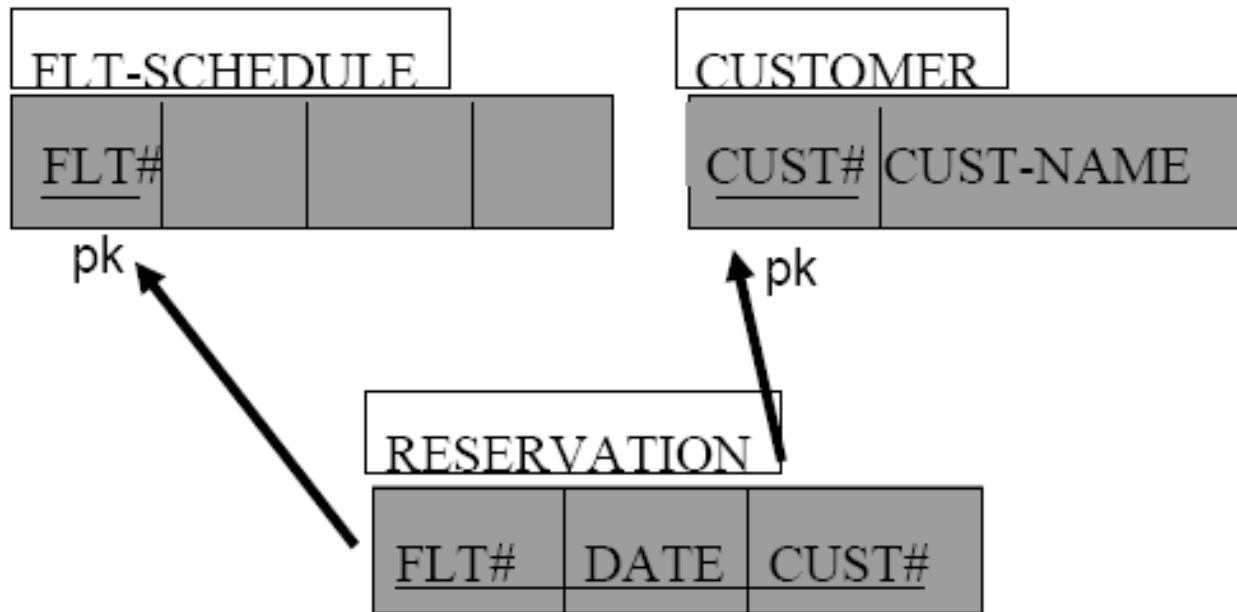
CUSTOMER

| | | | | | | | | |
|--------------|-------|--------|------|--------|--------|------|-------|-----|
| <u>cust#</u> | first | middle | last | phone# | street | city | state | zip |
|--------------|-------|--------|------|--------|--------|------|-------|-----|

RESERVATION

| | | | | | |
|-------------|-------------|--------------|-------|-----------------|----------------|
| <u>flt#</u> | <u>date</u> | <u>cust#</u> | seat# | check-in-status | <u>ticket#</u> |
|-------------|-------------|--------------|-------|-----------------|----------------|

Example 2 (Cont.)



Other Types of Constraints

- Domain constraints: values of each attribute A must be an atomic value from the domain for that attribute, $\text{dom}(A)$
- Semantic Integrity Constraints:
 - Based on application semantics and cannot be expressed in a database schema
 - E.g., “one student cannot register for more than 9 credits”
 - *A constraint specification language* may have to be used to express these constraints
 - SQL-99 uses triggers and assertions

Exercise

Consider the following relations for a database that keeps track of student enrollment in courses and the books adopted for each course:

STUDENT(SSN, Name, Major, Bdate)

COURSE(Course#, Cname, Dept)

ENROLL(SSN, Course#, Quarter, Grade)

BOOK_ADOPTION(Course#, Quarter, Book_ISBN)

TEXT(Book_ISBN, Book_Title, Publisher, Author)

Add foreign key constraints to the schema above

Update Operations on Relations

- INSERT a tuple
 - Domain, key, entity, and referential integrity constraints can be violated
- DELETE a tuple
 - Only referential integrity constraint can be violated
- MODIFY a tuple
 - Modify PK: delete one tuple and insert another in its place
 - Modify FK: make sure the new value refers to an existing tuple in the referenced relation (or is null)
 - Neither PK nor FK: only check new value is of the correct domain and data type

Update Operations on Relations (Cont.)

- Integrity constraints should not be violated by the update operations
- Several update operations may have to be grouped together
- Updates may *propagate* to cause other updates automatically

Update Operations on Relations (Cont.)

- In case of integrity violation, several actions can be taken:
 - Cancel the operation that causes the violation (REJECT option)
 - Perform the operation but inform the user of the violation
 - Trigger additional updates so the violation is corrected (CASCADE option, SET NULL option)
 - Execute a user-specified error-correction routine

Example

- The instances below satisfy some constraints

Climbers:

| CId | CName | Skill | Age |
|-----|---------|-------|-----|
| 123 | Edmund | EXP | 80 |
| 214 | Arnold | BEG | 25 |
| 313 | Bridget | EXP | 33 |
| 212 | James | MED | 27 |

Climbs:

| CId | RId | Date | Duration |
|-----|-----|----------|----------|
| 123 | 1 | 10/10/88 | 5 |
| 123 | 3 | 11/08/87 | 1 |
| 313 | 1 | 12/08/89 | 5 |
| 214 | 2 | 08/07/92 | 2 |
| 313 | 1 | 06/07/94 | 3 |

- Insert (123, Jeremy, MED, 16) into Climbers?
- Insert (456, 2, 09/13/98, 3) into Climbs?
- Delete (313, Bridget, EXP, 33) from Climbers?
- Modify 123 to 456 in Climbers?