

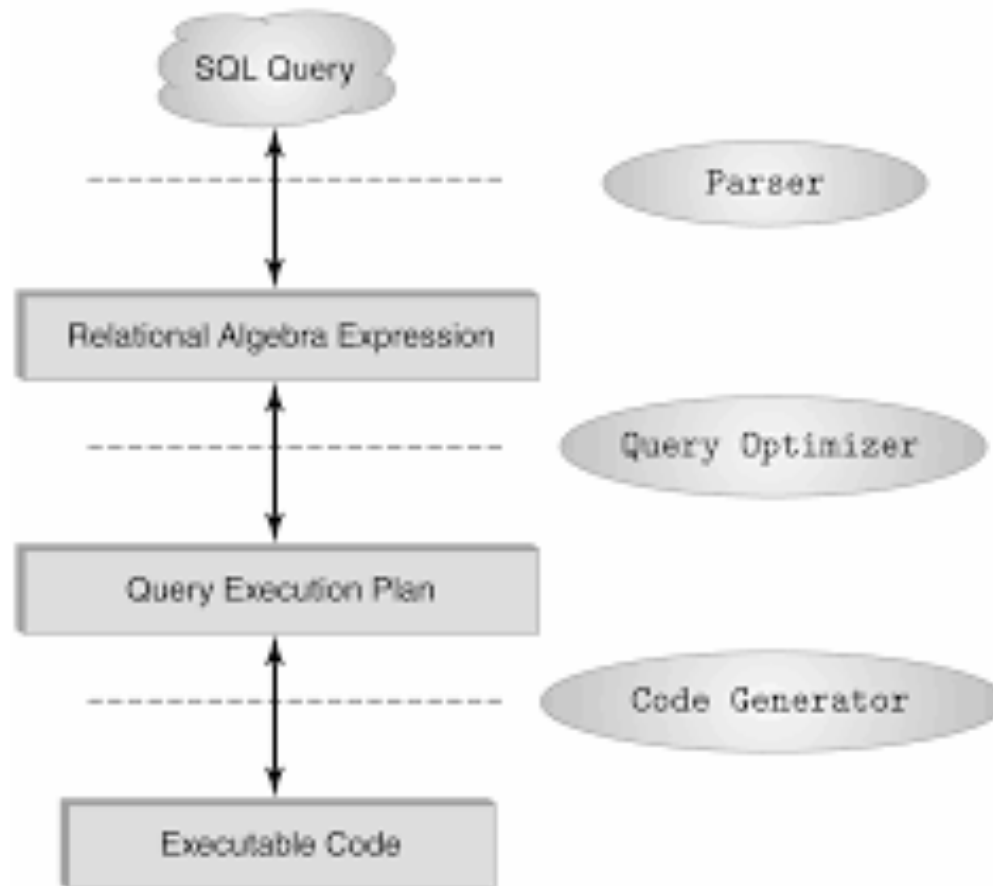
Lecture 6

Relational Algebra

What is Relational Algebra?

- A set of operations each of which takes a relation (or relations) as input and produces a relation as output
- **Procedural:** specify how to construct the result
- Two groups of operations
 - Set operations from mathematical set theory
 - Union
 - Intersection
 - Set Difference
 - Cartesian Product
 - Operations developed specially for relational databases
 - Project
 - Select
 - Join and others

Schematic View of Query Processing



SELECT

- Selects the tuples (rows) from a relation R that satisfy a certain *selection condition* c
- Form of the operation: $\sigma_c(R)$
- The condition c is a Boolean expression on the attributes of R of two forms:
 - $\langle \text{attribute name} \rangle \langle \text{comparison op} \rangle \langle \text{constant value} \rangle$
 - $\langle \text{attribute name} \rangle \langle \text{comparison op} \rangle \langle \text{attribute name} \rangle$
- Resulting relation has the *same attributes* as R and includes each tuple whose attribute values satisfy the condition c

SELECT Example

Routes

RIId	RName	Grade	Rating	Height
1	Last Tango	II	12	100
2	Garden Path	I	2	60
3	The Sluice	I	8	60
4	Picnic	III	3	400

$\sigma_{Height=60}$ Routes :

RIId	RName	Grade	Rating	Height
2	Garden Path	I	2	60
3	The Sluice	I	8	60

What Can Go in a Selection Condition?

- Conditions are built up from Boolean-valued operations on the field names
 - E.g., Height \geq 100, RName = 'Picnic'
- Clauses can be arbitrarily connected by the Boolean operator OR, AND, NOT to form a general selection condition
 - E.g., $\sigma_{(\text{Dno}=4 \text{ AND Salary} > 25000) \text{ OR } (\text{Dno}=5 \text{ AND Salary} > 30000)}(\text{EMPLOYEE})$

Characteristics of Selection

- Selection operator is unary
- Apply to a single relation R
- Test each tuple individually
- Degree of the relation resulting from a SELECT operation = the degree of R
- # of tuples in the resulting relation is always less than or equal to the # of tuples in R
- σ is commutative

PROJECT

- Select some columns from the table and discard other columns
- Keeps only certain attributes (columns) from a relation R specified in an *attribute list* L
- Form of operation: $\pi_L(R)$

PROJECT Example

Routes:

<u>RId</u>	<u>RName</u>	<u>Grade</u>	<u>Rating</u>	<u>Height</u>
1	Last Tango	II	12	100
2	Garden Path	I	2	60
3	The Sluice	I	8	60
4	Picnic	III	3	400

$\pi_{RId, Height}$ Routes :

<u>RId</u>	<u>Height</u>
1	100
2	60
3	60
4	400

Projection Discussion

- Suppose the result of a projection has a repeated value, how do we treat it?

π_{Height} Routes	<u>Height</u>	<u>Height</u>
	100	100
	60	60
	60	400
	400	

- In “pure” relational algebra the answer is always a set (the second answer)
- However SQL and some other languages return, by default, a multiset

Characteristics of Projection

- Projection operator is unary
- Apply to a single relation R
- The degree of the resulting relation is equal to the number of attributes in $\langle \text{attribute list} \rangle$
- # of tuples in the resulting relation is always less than or equal to # of tuples in R
- π is not commutative

Database Queries

- Queries are formed by building up expressions with the operations of the relational algebra
- For example, select-project expressions are very common:

$$\pi_{Name, Age}(\sigma_{Age \geq 30} \text{Climbers})$$

- What does this mean in English?
- Also, could we interchange the order of the σ and π ?
Can we always do this?

Sequences of Operations

- Write the operations as a single **relational algebra expression** by nesting the operations
- E.g., to retrieve the first name, last name, and salary of all employees who work in department number 5, we can apply a select and a project operation

$$\pi_{\text{FNAME, LNAME, SALARY}}(\sigma_{\text{DNO=5}}(\text{EMPLOYEE}))$$

- OR explicitly show the sequence of operations, giving a name to each intermediate relation:

$$\text{DEP5_EMPS} \leftarrow \sigma_{\text{DNO=5}}(\text{EMPLOYEE})$$

$$\text{RESULT} \leftarrow \pi_{\text{FNAME, LNAME, SALARY}}(\text{DEP5_EMPS})$$

Set Operations

- Binary operations from mathematical set theory
 - UNION: $R_1 \cup R_2$
 - INTERSECTION: $R_1 \cap R_2$
 - SET DIFFERENCE: $R_1 - R_2$
 - CARTESIAN PRODUCT: $R_1 \times R_2$

Union Compatibility

- For \cup , \cap , $-$, the operand relations $R_1(A_1, A_2, \dots, A_n)$ and $R_2(B_1, B_2, \dots, B_n)$ must be **union-compatible** if
 - R_1 and R_2 have the same number of attributes, and
 - $\text{dom}(A_i) = \text{dom}(B_i)$ for $i=1, 2, \dots, n$
- The resulting relation for $R_1 \cup R_2$, or $R_1 \cap R_2$, or $R_1 - R_2$ has the same attribute names as the *first* operand relation R_1 (by convention)



Do we need to check union compatibility for Cartesian product?


Yes

No




Do we need to check union compatibility for Cartesian product?

Yes




No






Do we need to check union compatibility for Cartesian product?

Yes



No



UNION

- Two relations are union-compatible
- $R_1 \cup R_2$: includes all tuples that are either in R_1 or in R_2 or in both R_1 and R_2
- Duplicates are eliminated

UNION Example

Climbers:

<u>CIId</u>	<u>CName</u>	<u>Skill</u>	<u>Age</u>
123	Edmund	EXP	80
214	Arnold	BEG	25
313	Bridget	EXP	33
212	James	MED	27

Hikers:

<u>CIId</u>	<u>CName</u>	<u>Skill</u>	<u>Age</u>
214	Arnold	BEG	25
898	Jane	MED	39

Climbers \cup Hikers :

<u>CIId</u>	<u>CName</u>	<u>Skill</u>	<u>Age</u>
123	Edmund	EXP	80
214	Arnold	BEG	25
313	Bridget	EXP	33
212	James	MED	27
898	Jane	MED	39

INTERSECTION

- Two relations are union-compatible
- $R_1 \cap R_2$: includes all tuples that are in both R_1 and R_2

INTERSECTION Example

Climbers:

<u>CId</u>	<u>CName</u>	<u>Skill</u>	<u>Age</u>
123	Edmund	EXP	80
214	Arnold	BEG	25
313	Bridget	EXP	33
212	James	MED	27

Hikers:

<u>CId</u>	<u>CName</u>	<u>Skill</u>	<u>Age</u>
214	Arnold	BEG	25
898	Jane	MED	39

Climbers \cap Hikers :

<u>CId</u>	<u>CName</u>	<u>Skill</u>	<u>Age</u>
214	Arnold	BEG	25

SET DIFFERENCE (or MINUS)

- Two relations are union-compatible
- $R_1 - R_2$: includes all tuples that are in R_1 but not in R_2

SET DIFFERENCE Example

Beginners:

<u>CId</u>	<u>CName</u>	<u>Skill</u>	<u>Age</u>
214	Arnold	BEG	25
212	James	MED	27

Climbers - Beginners:

<u>CId</u>	<u>CName</u>	<u>Skill</u>	<u>Age</u>
123	Edmund	EXP	80
313	Bridget	EXP	33

Climbers:

<u>CId</u>	<u>CName</u>	<u>Skill</u>	<u>Age</u>
123	Edmund	EXP	80
214	Arnold	BEG	25
313	Bridget	EXP	33
212	James	MED	27

Characteristics of Set Operations

- Both union and intersection are *commutative operations*

$$\mathbf{R \cup S = S \cup R, \text{ and } R \cap S = S \cap R}$$

- Both union and intersection are *associative operations*

$$\mathbf{R \cup (S \cup T) = (R \cup S) \cup T, \text{ and}}$$

$$\mathbf{(R \cap S) \cap T = R \cap (S \cap T)}$$

- The minus operation is *not commutative*

$$\mathbf{R - S \neq S - R}$$

Exercise

- Customer: C_Name, C_Street, C_City
- Depositor: C_Name, Account#
- Loan: B_Name, Loan_#, Amount
- Borrower: C_Name, Loan#

Exercise (Cont.)

- Select tuples of the loan relation where the branch name is “GMU”.
- Find all tuples in which the amount of loan is more than \$12000.
- Find those tuples pertaining to loans of more than \$12000 made by “GMU” branch.
- List all loan numbers and the amount of loan.
- Find the names of all customers who live in “Fairfax”.
- Find the names of all bank customers who have either an account or a loan or both.
- Find the names of all customers of the bank who have an account but not a loan.
- Find the names of all customers who have both a loan and an account.